Qwen2.5VL

Qwen2.5VL
- 1. Model Size
- 2. Performance
- 3. Using Qwen2.5VL
  - 3.1 Running Qwen2.5VL
  - 3.2 Start a conversation
  - 3.3 Visual Function
  - 3.4 Ending the Conversation
  - 3.5 Chinese Dialogue
- References

Note: Due to performance limitations, this model cannot be run on a Raspberry Pi 5 2GB. The 4GB version will run much slower.

Qwen2.5-VL is Qwen's new flagship visual language model and a significant leap forward compared to the previous Qwen2-VL.

1. Model Size

Model	Volume
qwen2.5vl:3b	3.2GB
qwen2.5vl:7b	6.0GB

2. Performance

		Qwen2.5-VL-3B	InternVL2.5-4B	Qwen2-VL-7B
College-level Problems	MMMU	53.1	52.3	**54.1**
	MMMU Pro	**31.6**	32.7	30.5
Document and Diagrams Reading	DocVQA	93.9	91.6	**94.5**
	InfoVQA	**77.1**	72.1	76.5
	CC-OCR	**74.5**	-	61.6
	OCRBenchV2	**54.3**	-	47.8
General Visual Question Answering	MegaBench	28.9	-	**34.4**
	MMStar	55.8	58.3	**60.7**
	MMBench1.1	77.6	79.3	**80.7**
Math	MathVista	**62.3**	60.5	58.2
	MathVision	**21.2**	20.9	16.3
Video Understanding	VideoMME	61.5	62.3	**63.3**
	MMBench-Video	1.6	**1.7**	1.4
	LVBench	**43.3**	-	-
	CharadesSTA	**38.8**	-	-
Visual Agent	AITZ	**76.9**	-	-
	Android Control	**63.7**	-	-
	ScreenSpot	**55.5**	-	55.3
	ScreenSpot Pro	**23.9**	-	1.6

		Qwen2.5-VL 7B	Qwen2-VL 7B	GPT-4o Mini	Other Best Open LVLM With Similar Size
College-level Problems	MMMU	58.6	54.1	**60.0**	56.0
	MMMU Pro	**38.3**	30.5	37.6	34.3
Document and Diagrams Reading	DocVQA	**95.7**	94.5	-	93.0
	InfoVQA	**82.6**	76.5	-	77.6
	CC-OCR	**77.8**	61.6	-	61.6
	OCRBenchV2	**56.3**	47.8	43	47.8
General Visual Question Answering	MegaBench	36.8	34.4	**43.1**	34.4
	MMStar	**63.9**	60.7	54.8	62.8
	MMBench1.1	**82.7**	80.7	76.0	79.4
Math	MathVista	**68.2**	58.2	52.5	67.2
	MathVision	**25.1**	16.3	-	19.7
Video Understanding	VideoMME	**65.1**	63.3	64.8	63.3
	MMBench-Video	**1.8**	1.4	-	1.7
	LVBench	**45.3**	-	-	38.4
	CharadesSTA	43.6	-	-	**48.4**
Visual Agent	AITZ	**81.9**	-	-	53.3
	Android Control	**60.1**	-	-	**61.5**
	ScreenSpot	84.7	55.3	-	**89.5**
	ScreenSpot Pro	29.0	1.6	-	**35.7**

3. Using Qwen2.5VL

3.1 Running Qwen2.5VL

Use the run command to start running the model. If the model is not already downloaded, it will automatically pull the model from the Ollama model library:

```
ollama run qwen2.5vl:3b
```

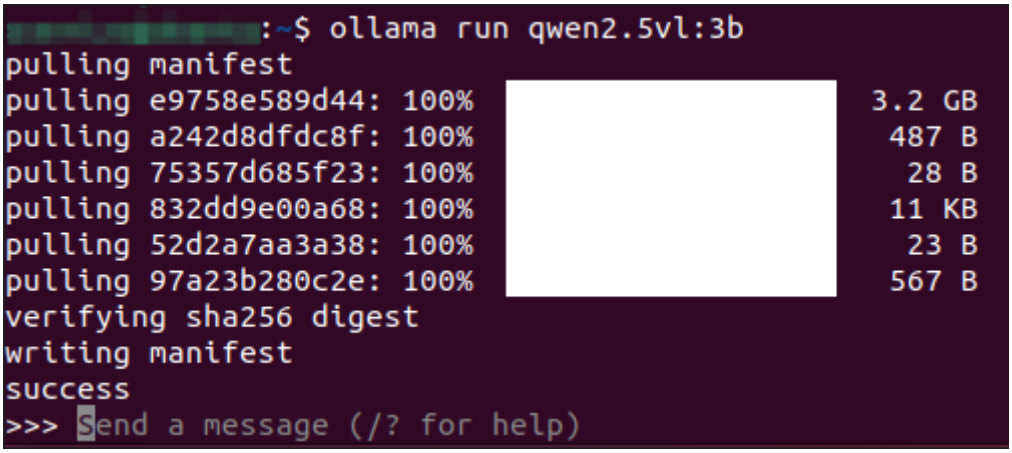

3.2 Start a conversation

```
Please tell me how many hours there are in a day
```

Response time depends on your hardware configuration. Please be patient!

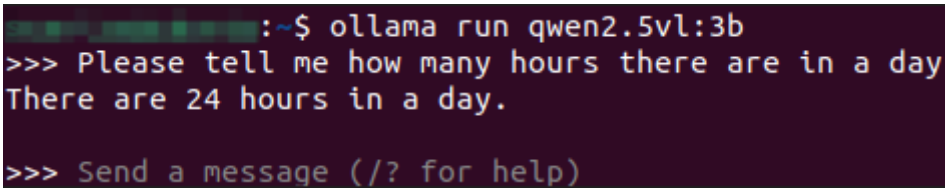

3.3 Visual Function

```
What do you see in this picture? :./test_pic.png
#Use ": + image path" in the conversation to enable the model to use its visual
function and interpret the information in the image.
```

```
:~$ ollama run qwen2.5vl:3b
>>> What do you see in this picture?:./test_pic.png
Added image './test_pic.png'
The picture shows a person sitting on a rooftop or elevated platform,
looking up towards the sky. The person is wearing a white tank top and
appears to be in a relaxed or contemplative pose. In the foreground, there
is a clear plastic bottle. The background features a bright, sunny sky
with some clouds and a cityscape with buildings visible in the distance.
The overall atmosphere of the image is calm and serene.
```

3.4 Ending the Conversation

Use the `Ctrl+d` shortcut or `/bye` to end the conversation!

3.5 Chinese Dialogue

If you don't have a Chinese input method, please refer to the Chinese input method switching tutorial.

Chinese Dialogue:

```
>>> 一个小时是多少分钟？请用中文回答我
一个小时是60分钟。

>>> Send a message (/? for help)
```

References

> **Ollama**

Official website:https://ollama.com/

GitHub：https://github.com/ollama/ollama

> **Qwen2.5VL**

GitHub：https://github.com/QwenLM/Qwen2.5-VL

Ollama corresponding model:https://ollama.com/library/qwen2.5vl